

# Homework 2

Abigail Wright

STA-4364

$$f(x_1, x_2) = 2x_1^2 + 2x_1x_2 + x_2^2 + x_1 - x_2$$

## Problem 1

1.

$$f(x_1, x_2) = \frac{1}{2} \begin{pmatrix} x_1 & x_2 \end{pmatrix} \begin{bmatrix} Q_{11} & Q_{12} \\ Q_{21} & Q_{22} \end{bmatrix} \begin{pmatrix} x_1 \\ x_2 \end{pmatrix} + \begin{pmatrix} b_1 & b_2 \end{pmatrix} \begin{pmatrix} x_1 \\ x_2 \end{pmatrix}$$

Expanding

$$f(x_1, x_2) = \frac{1}{2} \begin{pmatrix} x_1 & x_2 \end{pmatrix} \begin{bmatrix} Q_{11} & Q_{12} \\ Q_{21} & Q_{22} \end{bmatrix} \begin{pmatrix} x_1 \\ x_2 \end{pmatrix} + \begin{pmatrix} b_1 & b_2 \end{pmatrix} \begin{pmatrix} x_1 \\ x_2 \end{pmatrix}$$

$$f(x_1, x_2) = \frac{1}{2} (Q_{11}x_1^2 + 2Q_{12}x_1x_2 + Q_{22}x_2^2) + b_1x_1 - b_2x_2$$

Thus:

$$2x_1^2 + 2x_1x_2 + x_2^2 + x_1 - x_2$$

$$Q_{11} = 4$$

$$Q_{12} = 1$$

$$Q_{22} = 2$$

Thus,

$$Q = \begin{bmatrix} 4 & 1 \\ 1 & 2 \end{bmatrix}$$

$$b = \begin{pmatrix} 1 \\ -1 \end{pmatrix}$$

Defined as:

```
def compute_f():
    x = np.array([1, 1])
    x_T = np.array([[1],
                    [1]])

    Q11 = 4
    Q12 = 1
    Q21 = 1
    Q22 = 2
    Q = np.array([[Q11, Q12], [Q21, Q22]])

    b1 = 1
    b2 = -1
    b = np.array([b1, b2])
    return Q, b
```

## 2.

The gradient of  $f(x_1, x_2)$  is given by:  $\nabla f = Qx + b$

The first derivative with respect to  $x_1$ :

$$f'(x_1, x_2) = \frac{d}{dx_1}(2x_1^2 + 2x_1x_2 + x_2^2 + x_1 - x_2) = 4x_1 + 2x_2 + 1 \quad (1)$$

The first derivative with respect to  $x_2$ :

$$f'(x_1, x_2) = \frac{d}{dx_2}(2x_1^2 + 2x_1x_2 + x_2^2 + x_1 - x_2) = 2x_1 + 2x_2 - 1 \quad (2)$$

Where  $Qx + b = \begin{bmatrix} 4x_1 + 2x_2 + 1 \\ 2x_1 + 2x_2 - 1 \end{bmatrix}$

Utilizing the code:

```
def f(x1,x2):
    return 2 * x1**2 + 2 * x1 * x2 + x2**2 + x1 - x2

def gradient_f(x1,x2):
    return np.array([4 * x1 + 2 * x2 + 1, 2 * x1 + 2 * x2 -1])
```

### 3.

The formula for the gradient requires us to define our original function and first derivative of the function. Which is defined in our code as:

```
def f(x1,x2):  
    return 2 * x1**2 + 2 * x1 * x2 + x2**2 + x1 - x2  
  
def gradient_f(x1,x2):  
    return np.array([4 * x1 + 2 * x2 + 1, 2 * x1 + 2 * x2 - 1])
```

Where  $f(x_1, x_2) = 2x_1^2 + 2x_1x_2 + x_2^2 + x_1 - x_2$  and  $\nabla(f(x_1, x_2)) = \begin{bmatrix} 4x_1 + 2x_2 + 1 \\ 2x_1 + 2x_2 - 1 \end{bmatrix}$

### 4.

The steps used for the Steepest Descent are as follows:

- Set counter  $k = 0$  iterations.
- Compute  $\nabla f(x_k)$ .
- Choose  $\lambda_k$ .
- Update  $x_{k+1} = x_k - \lambda_k \nabla f(x_k)$ .
- Continue until  $\|\nabla f(x_k)\| < \text{tolerance}$ .

To find the minimizer of  $f$  using gradient descent with exact line search, we define our function for each iteration by:

```
def g(lambda_k, x, r):  
    return f(x - lambda_k * r)
```

We also define our steepest descent function by:

```
def exact_line(f, gradient_f, x0, tol=1.e-8, maxit=100):  
    x = np.array(x0)  
    r = gradient_f(x)  
    k = 0  
    while np.abs(np.linalg.norm(r)) > tol and k < maxit:
```

```

        lambda_k = spo.golden(lambda l: g(l, x, r))
        x = x - lambda_k * r
        r = gradient_f(x)
        k += 1
    return x, k

```

In this case I selected  $x_0 = [4.0, 2.0]$  which resulted in minimizer  $x : [-1.0, 1.49999999]$  with 19 iterations before convergence.

## 5.

Here when it comes to fixed step-size we will follow similar steps as the exact line search with some updates:

- Choose a set value for  $\lambda_k$ .
- Compute  $\nabla f(x_k)$ .
- Update  $x_{k+1} = x_k - \lambda_k \nabla f(x_k)$ .
- Continue until convergence is reached

The code used to find the minimizer of  $f$  using gradient descent with fixed step-size:

```

def g(lambda_k,x,r):
    return f(x - lambda_k*r)

def fixed_step(f, gradient_f, x0, tol=1.e-6, maxit=100):
    step = [x0]
    x = x0
    r = gradient_f(x)
    lambda_k = 0.2

    for i in range(maxit):
        diff = lambda_k * r
        if npl.norm(diff)<tol:
            break
        x = x - diff
        r = gradient_f(x)

```

```

        step.append(x) ## tracking
    return step, i + 1

```

In this case I selected  $x_0 = [2.0, 1.0]$  with  $\lambda = 0.2$  which resulted in minimizer  $x : [-0.99999997, 1.49999994]$  with 105 iterations before convergence.

## 6.

Find the minimizer of  $f$  using backtracking we must follow the steps:

- Set counter  $k = 0$  iterations.
- Make an initial guess for  $\mathbf{x}_0$
- Choose an initial  $\alpha = 1$ .
- Update  $\lambda_{k+1} = \beta\lambda - k$
- Go until  $f(x_{k+1} - \lambda_k \nabla f(x_k)) \leq f(x_{k+1} - \alpha \lambda_k \|\nabla f(x_k)\|^2)$
- Calculate  $x_{k+1} = x_k - \lambda_k \nabla f(x_k)$  and update  $k = k + 1$
- Continue until  $\|\nabla f(x_k) < \text{tolerance}\|$

The code used to find the minimizer of  $f$  using gradient descent with back-track line search included defining step-size, gradient descent, and backtracking functions:

```

def step_size(f, gradient_f, x):
    alpha = 1.0
    beta = 0.8
    r = gradient_f(x)
    while f(x - alpha*r) > (f(x) - 0.5*alpha*lp.norm(r)**2):
        alpha *= beta
    return alpha

def g(lambda_k,x,r):
    return f(x - lambda_k*r)

def back_track(f, gradient_f, x0, tol=1.e-8, maxit=100):
    x = np.array(x0)

```

```

r = gradient_f(x)
k = 0
while npl.norm(r) > tol and k < maxit:
    lambda_k = step_size(f, gradient_f, x)
    x = x - lambda_k * r
    r = gradient_f(x)
    k += 1
return x, k

```

In the case of backtracking line search, our convergence was reached much faster when compared to other methods. For  $x_0$  the point  $[2.0, 1.0]$  was selected with  $\alpha = 1$  and  $\beta = 0.08$ . Which resulted in minimizer  $x : [-1.0, 1.5]$  with 61 iterations before convergence.

## Problem 2

### 1.

First, we must find the first derivative with respect to  $\beta_0$  and  $\beta_1$ :

$$g(\beta_0, \beta_1) = \begin{bmatrix} \frac{\partial l}{\partial \beta_0} \\ \frac{\partial l}{\partial \beta_1} \end{bmatrix} \text{ where } y = \begin{bmatrix} 3.8 \\ 6.5 \\ 11.5 \end{bmatrix} \text{ and } x = \begin{bmatrix} 1 & 5 \\ 1 & 6 \\ 1 & 7 \end{bmatrix}$$

The derivative with respect to  $\beta_0$  is:  $\frac{\partial l}{\partial \beta_0}$

$$= -\frac{1}{3} [(3.8 - \beta_0 - 5\beta_1) + (6.5 - \beta_0 - 6\beta_1) + (11.5 - \beta_0 - 7\beta_1)]$$

The derivative with respect to  $\beta_1$  is:  $\frac{\partial l}{\partial \beta_1}$

$$= -\frac{1}{3} [(5)(3.8 - \beta_0 - 5\beta_1) + (6)(6.5 - \beta_0 - 6\beta_1) + (7)(11.5 - \beta_0 - 7\beta_1)]$$

The array of the y-intercept and slope is defined as:  $\beta = [\beta_0 \quad \beta_1]^T$

The gradient is:  $g(\beta) = \frac{1}{N} x^T e$

Here we must create a function for  $g(\beta)$ :

```
def grad_beta(x, y, B_0, B_1):
```

```

n = len(y)
residuals = (y - (B_0 + (B_1 * x)))

df_yint = (1/n) * np.sum(residuals)
df_slope = (1/n) * np.sum((residuals) * x)
g_B = np.array([[df_yint],
                 [df_slope]])
return g_B

```

Here we see the gradient is  $[7.26666667, 46.16666667]$ .

## 2.

To find the minimizer of  $L$  using gradient descent with a fixed step-size:

- $g(\beta_0, \beta_1) = \begin{bmatrix} \frac{dl}{d\beta_0} \\ \frac{dl}{d\beta_1} \end{bmatrix}$
- $\beta = (\beta_0, \beta_1)^T$
- $\beta^{k+1} = \beta^{(k)} - \lambda \nabla g(\beta^k)$

Here using fixed step-size we are choosing a constant value for  $\lambda$  in order to move in direction of the gradient = 0.

Utilizing the defined code:

```

def gradient_descent_fixed_step(x, y, B_0_init,
                                B_1_init, lambda_k, tol=1e-8, maxit=200):
    B = np.array([B_0_init, B_1_init])
    steps = [B.copy()]
    k = 0

    for _ in range(maxit):
        gradient = grad_beta(x, y, B[0], B[1])
        B_new = B - lambda_k * gradient

        steps.append(B_new.copy())

    if np.linalg.norm(gradient) < tol:

```

```

        break

    B = B_new
    k += 1

return B, k

```

### 3.

Here we will define the backtracking function by:

```

def gradient_descent_backtracking(x, y, B_0_init,
    B_1_init, tol=1e-8, maxit=200):
    B = np.array([B_0_init, B_1_init])
    step = [B.copy()]
    k = 0

    for _ in range(maxit):
        gradient = grad_beta(x, y, B[0], B[1])
        lambda_k = backtracking_line_search(x, y, B, gradient)
        B_new = B - lambda_k * gradient

        step.append(B_new.copy())

        if np.linalg.norm(gradient) < tol:
            break

    B = B_new
    k += 1

return B, k

```

## Problem 3

### 1.

Write a function for the gradient  $g(\beta)$

Here we are going to find the  $g(\beta)$  by utilizing matrix multiplication:



Where  $y = \begin{bmatrix} 3.8 \\ 6.5 \\ 11.5 \end{bmatrix}$ ,  $x = \begin{bmatrix} 1 & 5 \\ 1 & 6 \\ 1 & 7 \end{bmatrix}$ , and  $\beta = [\beta_0 \quad \beta_1]^T$

To find the gradient using matrix multiplication:  $g(\beta) = \frac{1}{n} \mathbf{X}^T \mathbf{e}$

Utilizing  $\mathbf{e} = \mathbf{y} - \mathbf{x}\beta$  Here we must create a function for  $g(\beta)$ :

```
def grad_beta(x, y, B):
    B = B_0, B_1
    n = len(y)
    e = y - np.dot(x, B)
    g_B = (1/n) * np.dot(x.T, e)
    return g_B
```

Here we see the gradient is  $[7.26666667, 46.16666667]$ .

## 2.

To find the minimizer of  $L$  using gradient descent with fixed step-size with matrix multiplication:

- $g(\beta_0, \beta_1) = \begin{bmatrix} \frac{dl}{d\beta_0} \\ \frac{dl}{d\beta_1} \end{bmatrix}$
- $\beta = (\beta_0, \beta_1)^T$
- $e = y - x\beta$
- $g(\beta) = \frac{1}{n} x^T e$

The code for fixed step-size with matrix multiplication:

```
def gradient_descent_fixed_step(x, y, B_0_init, B_1_init,
                                lambda_k, tol=1.e-8, maxit=10000):
    B = np.array([B_0_init, B_1_init])
    steps = [B.copy()]
```

```

losses = [loss_function(x, y, B[0], B[1])]
k = 0

for _ in range(maxit):
    gradient = grad_beta(x, y, B[0], B[1])
    B_new = B - lambda_k * gradient

    steps.append(B_new.copy())
    loss = loss_function(x, y, B_new[0], B_new[1])
    losses.append(loss)

    if np.abs(losses[-1] - losses[-2]) < tol:
        break

    B = B_new
    k += 1

return steps, losses, k

```

Resulting in  $\beta_0 = -15.831477$  and  $\beta_1 = 3.849696$ . Which was reached after 17058 iterations utilizing fixed step size.

### 3.

To find the minimizer of  $L$  using gradient descent with backtracking we will choose an  $\alpha$  between 0.01 and 0.3 and a  $\beta$  between 0.1 and 0.3.

To define our required functions for backtracking is as follows:

```

def backtracking_line_search(x, y, B, gradient, alpha=0.5, beta=0.8):
    t = 1.0
    while np.linalg.norm(grad_beta(x, y, B - t * gradient))
    > (1 - alpha * t) * np.linalg.norm(gradient):
        t *= beta
    return t

def gradient_descent_backtracking(x, y, B_init, tol=1e-6, max_iter=10000):
    B = B_init.copy()
    steps = [B.copy()]

```

```

for i in range(max_iter):
    gradient = grad_beta(x, y, B)
    step_size = backtracking_line_search(x, y, B, gradient)
    B_new = B - step_size * gradient
    steps.append(B_new.copy())

    if np.linalg.norm(B_new - B) < tol:
        break

    B = B_new

return steps, i + 1

```

Resulting in  $\beta_0 = -0.005954$  and  $\beta_1 = 1.258902$ . Which was reached after 18 iterations utilizing backtracking exact line search.